

Networking

Review

What is the point of RTTI?

- a) It lets you discover the properties of classes you wrote
- b) It lets you discover the properties of classes you downloaded
- c) It lets you discover the properties of system classes
- d) It lets you discover the properties of parent classes

Review

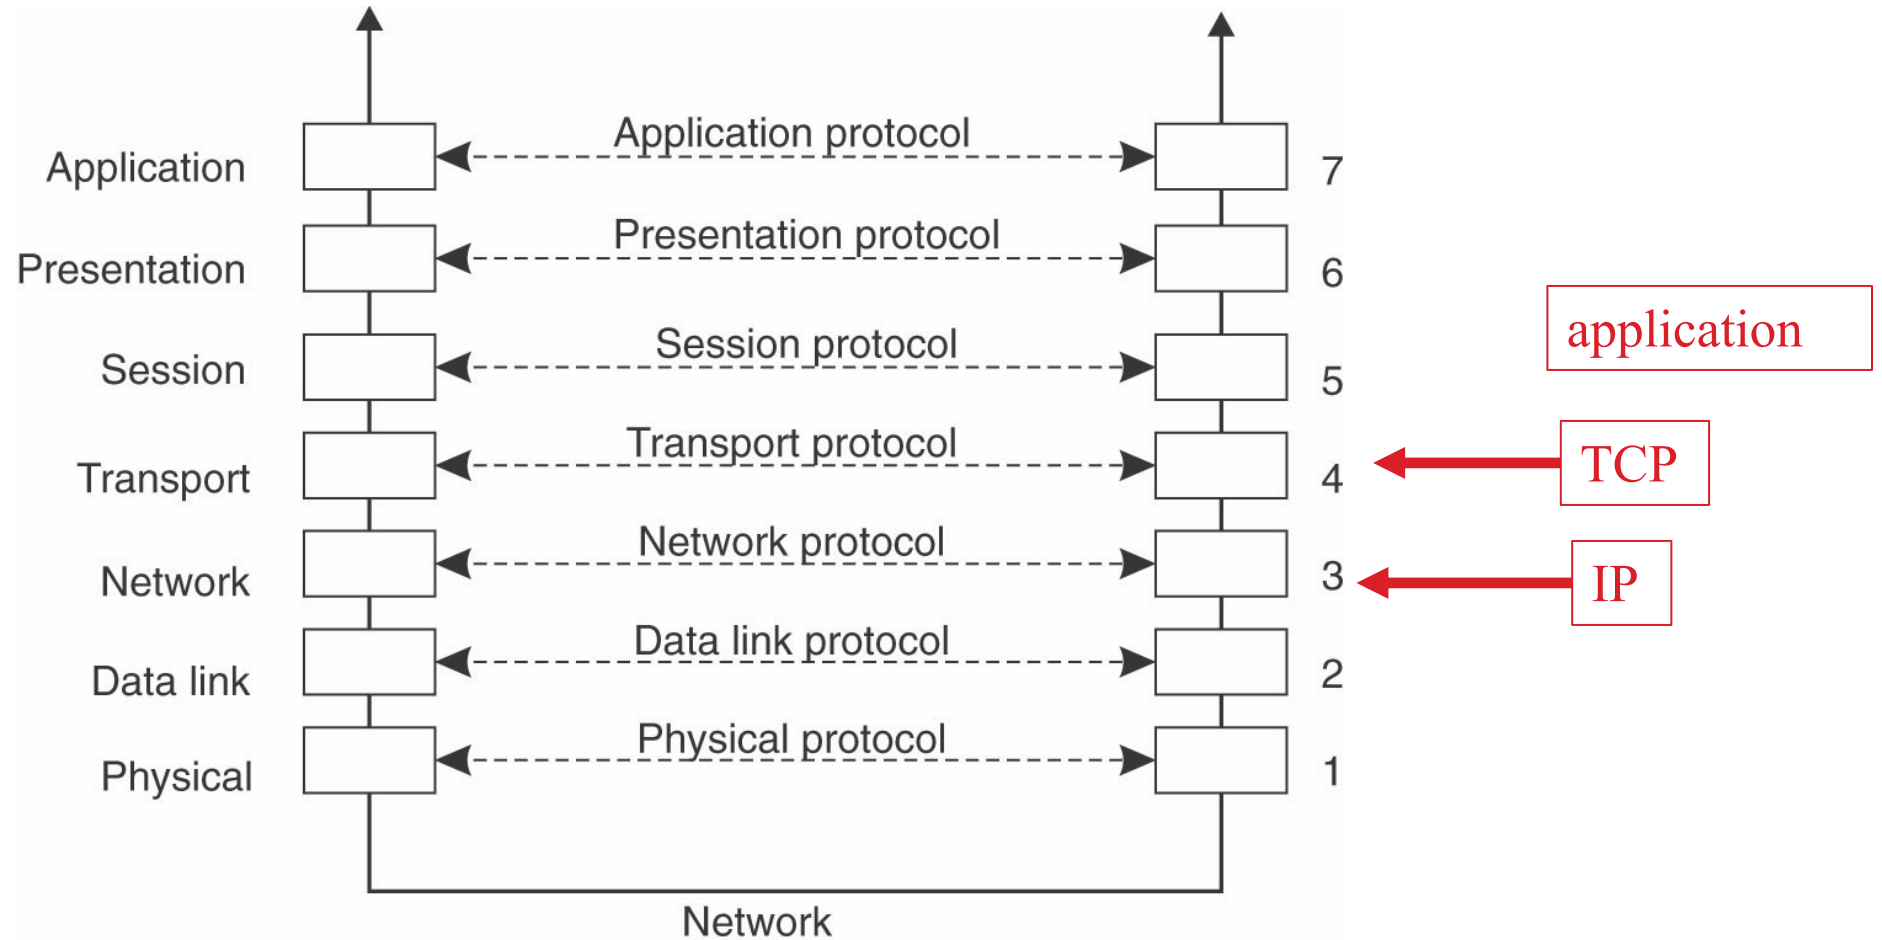
Which of these can be a problem with Generics?

- a) They are extremely inefficient
- b) They only work when also using interfaces
- c) Template method calls may not work
- d) They allow generic member data, but not generic methods

Objectives

- After this lecture, you will be able to:
 - define basic networking concepts
 - write Java programs that use sockets
 - understand the basics of the Java RMI protocol

OSI Model, 7 Layers



Tanenbaum & Van Steen

Low Level Communication

- OSI's upper levels are not always used
- Three “real” levels: application, connection, physical
- Connection level: two levels in OSI: transport and network
 - Network layer: *IP (Internet Protocol)*
 - host-to-host, connectionless
 - Transport layer: either
 - *UDP (User Datagram Protocol)*: process-to-process, connectionless, unreliable: fire and forget (and hope it gets there)
 - *TCP (Transmission Control Protocol)*: process-to-process, structured, connection-oriented, reliable via ack/nack+retransmit; re-orders subpackets on arrival: if things don't arrive correctly, re-do

TCP/IP

- These are usually grouped as *TCP/IP*: Transport Control Protocol and Internet Protocol (even though UDP is in there, too)
 - They rely on even lower-level protocols, down to the actual hardware
- Other, higher-level protocols or applications use their services
 - sendmail, ssh, http, ftp, ...
 - socket programming
 - Remote Method Invocation (RMI)

Networking Concepts, cont.

- Basic ideas:
 - **DNS** (Domain Name Service) address: four-part numeric address of a computer; it usually has a multipart string alias. Examples:
192.168.10.100; cmu.heinz.edu
 - **URL** (Uniform Resource Locator): network name. Example:
<http://cmu.heinz.edu/media/article1.html>
 - A URL must be translated (looked up, resolved) to a numeric DNS address

Networking Concepts, cont.

- ***Socket***: software record (class/object) for accessing the network
- ***Port***: integer connection number from 0-65535. These are logical, not physical, id's. 0-1023 are reserved (e.g. 22 is for ssh); 1024-49151 are for apps that may or may not be registered; the rest are for your program
 - If a port is in use, you can try another one
 - For security purposes, ports may be closed, or made not accessible through a firewall

https://en.wikipedia.org/wiki/List_of_TCP_and_UDP_port_numbers

Networking Concepts, cont.

- Think of the socket like a *street address*: your mail gets sent here.
- Then the port is like your *apartment number*: not all mail sent to this street address belongs to you; you should only receive the stuff with your apartment number on it

Networking Concepts, cont.

- Typically, one application is listening to one port on a server; some of those are system applications (like a mail server or a web server); many ports are not used and are available for an app that you write – with the restrictions mentioned above.
- But some apps need multiple listening ports, that's okay, just more complicated

Networking Concepts, cont.

- ***Client***: app that connects to a server to use its services by first making a connection request, then sending a message and (optionally) getting a reply
- ***Server***: app that sets up a socket and waits for client connection requests, get service requests, and (optionally) replies
 - Server is started first - because it waits for a connection request, it just pauses.
 - Client (eventually) asks to connect, then messages can be exchanged

Networking Concepts, cont.

- How client and server interact depends on the app
 - *Web server*: client sends URL, server sends HTML-encoded web page (or something more complicated)
 - *Database server*: client sends SQL queries, server sends tuples and/or acknowledgement (e.g. INSERT doesn't return tuples)
 - *Chat server*: exchange of String messages
 - *Peer-to-peer* networks: every machine is both a client and a server; data is either replicated across multiple machines or partitioned across multiple machines; finding things is more complicated
 - and many, many more

Sockets in Java

- The `Socket` and `ServerSocket` classes use the computer's network card
 - "Socket" should have been named "ClientSocket"
- The server knows (can look up) its own address
 - But it chooses a port to listen to
 - *Port stepper*: code that tries to connect using a port; if it fails, increments the port number, tries again
- The client doesn't care about its own address
 - It choose a server by name AND port number
 - Its (return) address is given to the server via TCP

Server Sockets

- `ServerSocket`: uses the port number the app will listen to
- Then call `accept ()` to wait for a client connection request
 - Blocking call: waits for the connection
 - Returns a `Socket` object
- Then read requests and write replies to the `Socket` object
 - Why doesn't it directly use the `ServerSocket` for that? So that it can listen for other clients
 - This should use a new `Thread` for each connection

Client Sockets

- `Socket`: use the server's address and port number that the app will connect to, then write to the socket and read answers back
 - in a simple app, the client hard-codes these or gets them from the user
 - in real apps, they might come from a broker (database of available servers) or name server
- The client **must** know the server's address and the port it's listening to
 - If you mess up the address or the port, your "mail" goes to the wrong place and is discarded
 - Unless you're using a lookup service

Outline of Server Code

```
// Server - missing some important details
ServerSocket s = new ServerSocket(portNumber);
Socket c = s.accept();
while (some condition) {
    message = get stuff from c's inputStream
    write a reply to c's outputStream
}
```



int

Server Code

```
ServerSocket serverSocket = null;
Socket clientConnection = null;
try {
    serverSocket = new ServerSocket(8001);
    clientConnection = serverSocket.accept();
    // Get the input connection to read from the client
    // Get the output connection to send replies to the client
    Scanner in = new Scanner(clientConnection.getInputStream());

    PrintWriter out = new PrintWriter(
        new BufferedWriter(
            new OutputStreamWriter(
                clientConnection.getOutputStream())));
```

for example: we're hoping this
port is available



input connection



output connection



Server Code, cont.

```
// Get a message from the input stream
String message = in.nextLine();
// Send a message on the output stream
out.println("OKAY");
} catch (IOException e) {
    e.printStackTrace();
}
```

Outline of Client Code

String URL

int

```
// Client - missing some important details
Socket s = new Socket(hostname, portNumber);
while (some condition) {
    write a message to s's outputStream
    message = get answer from s's inputStream
}
```

Client Code

```
Socket clientSocket = null;  
try {  
    clientSocket = new Socket("localhost", 8001);
```

for example

must match the
port the server is
listening to

```
    BufferedReader in = new BufferedReader(  
        new InputStreamReader(clientSocket.getInputStream()));
```

```
    PrintWriter out = new PrintWriter(  
        new BufferedWriter(  
            new OutputStreamWriter(clientSocket.getOutputStream()));
```

input connection

output connection

Client Code, cont.

```
// Send a message to the server
out.println("Hello from earth");
// Read the return message
String inMessage = in.readLine();
System.out.println("Client received back: " + inMessage);
} catch (IOException e) {
    e.printStackTrace();
}
```

Sockets, cont.

- For testing purposes:
 - "localhost" is used at the server's address in the `Socket()` constructor to connect to your own machine
 - Normally, use the real server address, and the programs are run on different machines
 - Run the server first, then run the client – two separate programs running on the same machine
 - Recall that the server will wait at the `accept()` call

Sockets, cont.

- Real servers run "forever" – that is, infinite loops - and are stopped manually
- Real servers start a Thread to handle the client Socket
 - Then they loop back to `accept()`, so multiple clients can be handled in parallel

Sockets, cont.

- You can send objects across the connection instead: use the `ObjectInputStream` and `ObjectOutputStream` classes instead
- Both the client and the server must have exactly the same class definition that you're reading/writing
- That class must implement `Serializable`

High Level Communication

- Programmers use function or procedure calls to structure code, so why not use that for communication?
 - ***Remote Procedure Call*** (RPC): mimic conventional function call – and – return, with parameters and return value
 - hides (most) low-level details of communication => transparency
 - ***Remote Method Invocation*** (RMI): object-oriented version of RPC, with some nicer features
- Java uses (can use) RMI

Java Remote Method Invocation

- Define an interface based on Remote:

```
public interface myInterface extends Remote
{
    returntype fn( params ) throws RemoteException;
}
```

Note this



- Server implements the interface, and inherits from UnicastRemoteObject

```
public class myClass extends UnicastRemoteObject implements
myInterface {
    ... code for fn( ) ...
}
```

and this



Java RMI, cont.

- Server generates stubs using *rmic*
- Server registers with registry service from `main()` using `Naming` class:
`bind(), rebind(), lookup()`
- Server started in background
 - "Real" servers run forever loops, so they have to be stopped manually
 - Alternatively, some special message can signal "stop"

Java RMI, cont.

- Client uses registry service to find server
- Client calls remote method

```
public class client {  
    public static void main(String[] args) {  
        String url = "rmi://someserver.com";  
        myclass m = (myclass)Naming.lookup(url + "binding name");  
        m.fn( ); // Calls remote method fn( )  
    }  
}
```

Java RMI, cont.

- Java also supports:
 - *callbacks*: the client registers with the server to respond to server-side events
 - *reflection*: request message can contain information about what method should be called, via `invoke()`