

Multithreading, part 2

Protecting your data

Review

Which is **not** a way to create threads in Java?

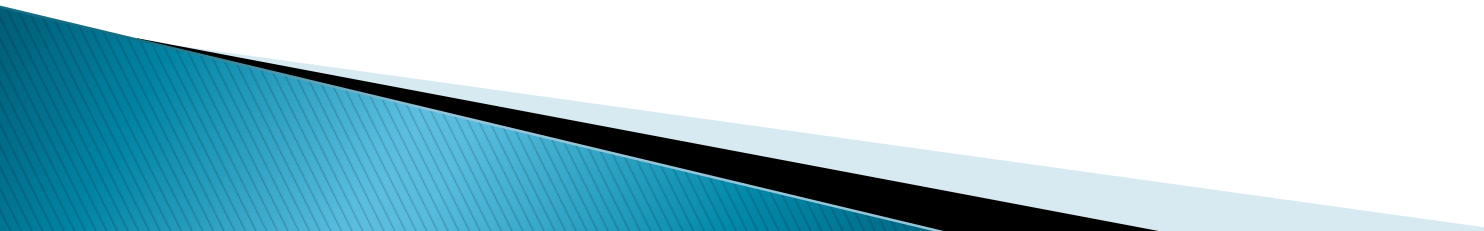
- a) Use a class that implements Runnable
- b) Use a class that extends Thread
- c) Use a class that extends Threadable
- d) Use an Executor

Review

What advantage does threading have?

- a) It always makes your program faster
- b) It never corrupts your data
- c) It's easy to debug
- d) It can improve performance, but it might corrupt your data and be hard to debug

Objectives

- Topics
 - Race conditions
 - Synchronization, locks, thread-safe data structures
 - Goals: after this lecture, you will be able to
 - understand why data shared by two or more threads can become corrupted
 - describe the basic ways to protect shared data
 - write threaded programs using thread-safe techniques
- 

Data Consistency

- If threads are independent – they work on separate data, like the subarrays above – then consistency is not a problem
 - More likely, they share data
- If data is shared by two or more threads, they must be careful updating the data. ***Race condition***: when shared data can become inconsistent (corrupted) due to updates from multiple threads.
- Even simple updates on primitive types (int, float, ...) can have race conditions

Consistency, cont.

- *Shared data* means that two or more threads use a *reference* to the same thing
 - Remember that "reference" here means pointer to an object on the heap, not primitive data
- This happens when the main program creates an object and gives it to (passes it as a parameter to an overloaded constructor of) a class that implements Runnable
- Even if you increment an int inside a shared class, it could go wrong – data gets corrupted

Consistency, cont.

- Race conditions result from bad timing (on the part of the thread scheduler, not your fault), running threads in parallel, and lack of data protection (your fault!)
- It stems from the fact that when you operate on data, it has to be fetched from main memory into the cpu/core, operated on there, and written back to memory
 - This takes a small but non-zero amount of time
 - *A thread can be interrupted in the middle of it* – by the operating system for some good reason; you have no control over it – that's the scheduler part

Example

- Suppose two threads *share* this: `int count = 0;` (inside an object, because they cannot share a primitive)
- Suppose both threads execute this instruction: `count++;`
- If the threads execute *serially* – one after the other – then there's no problem, and `count` ends up as 2
- How could such a simple operation go wrong?

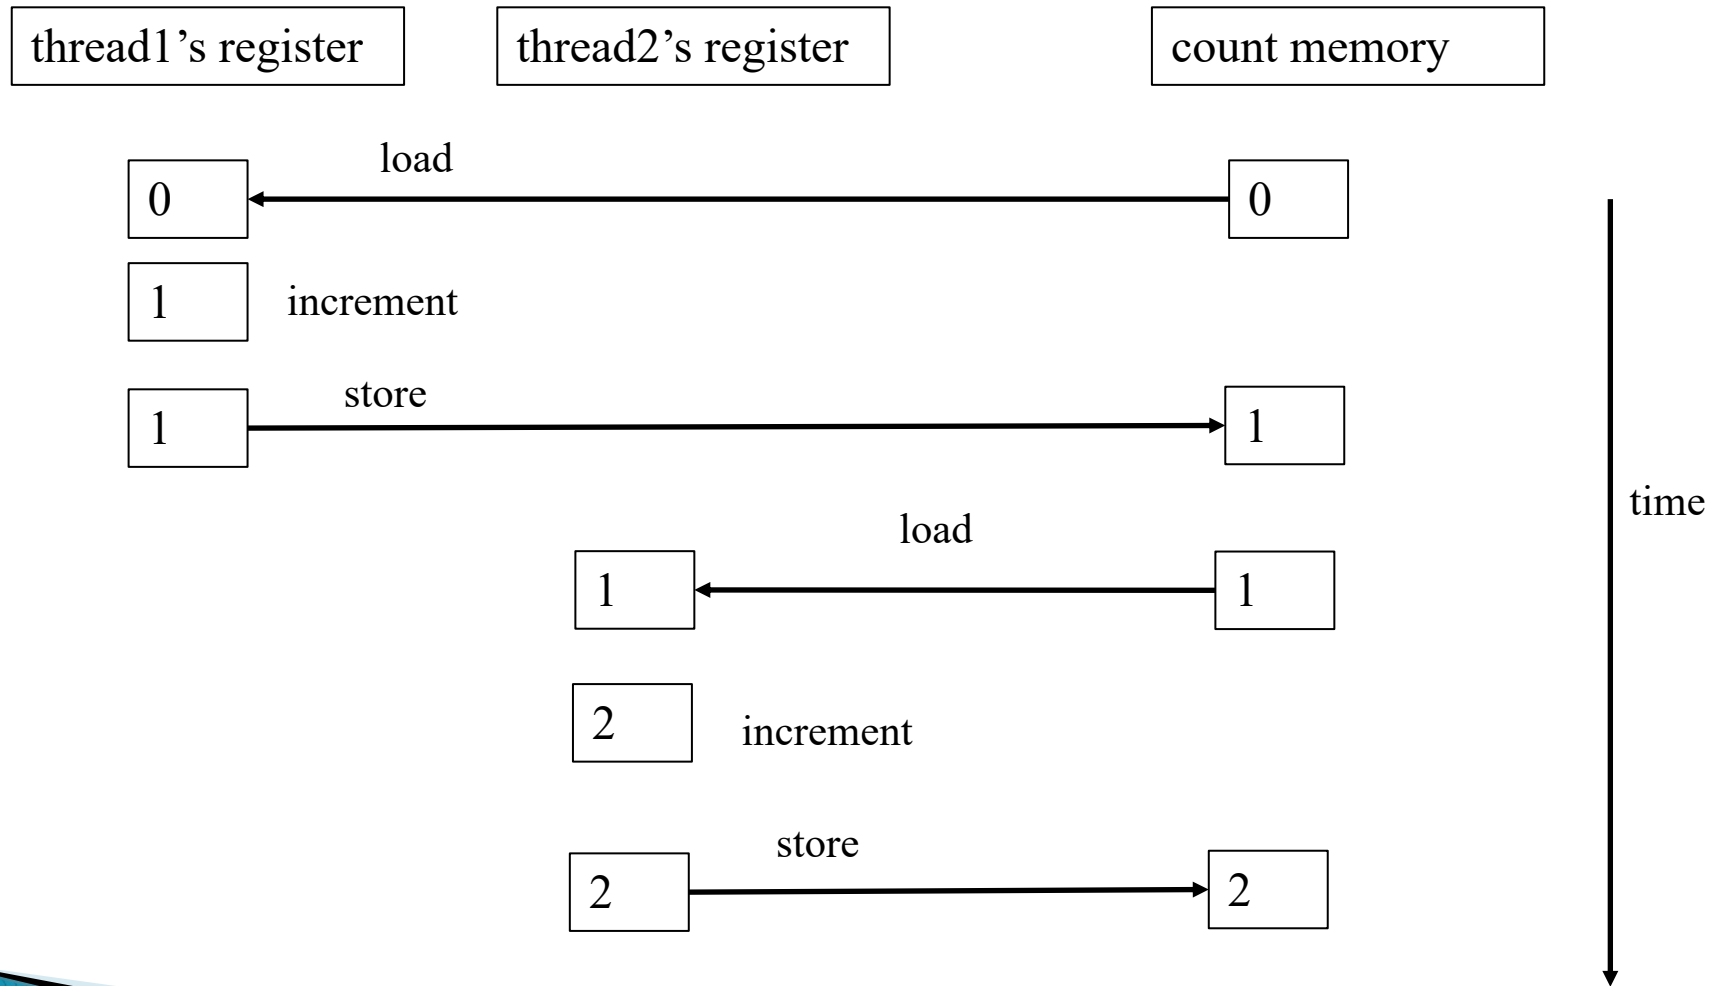
Example, cont.

- `count++` is not *atomic* (executes all at once or not at all) – because the bytecode looks something like this:

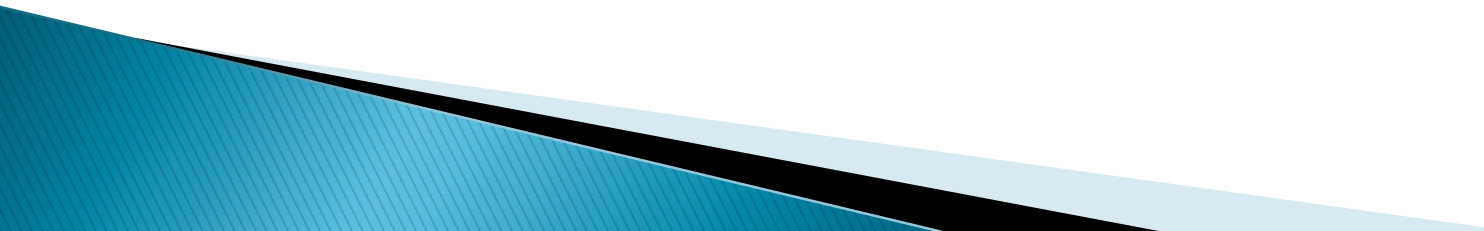
```
LOAD count;    // read count from main memory into the core
INCR count;    // add one into the core
STORE count;   // write count back to main memory
```

- Each core has local memory, called registers, separate from main memory
 - The number of registers is fairly small compared to the amount of main memory
 - You cannot operate on data in memory directly (well ...)

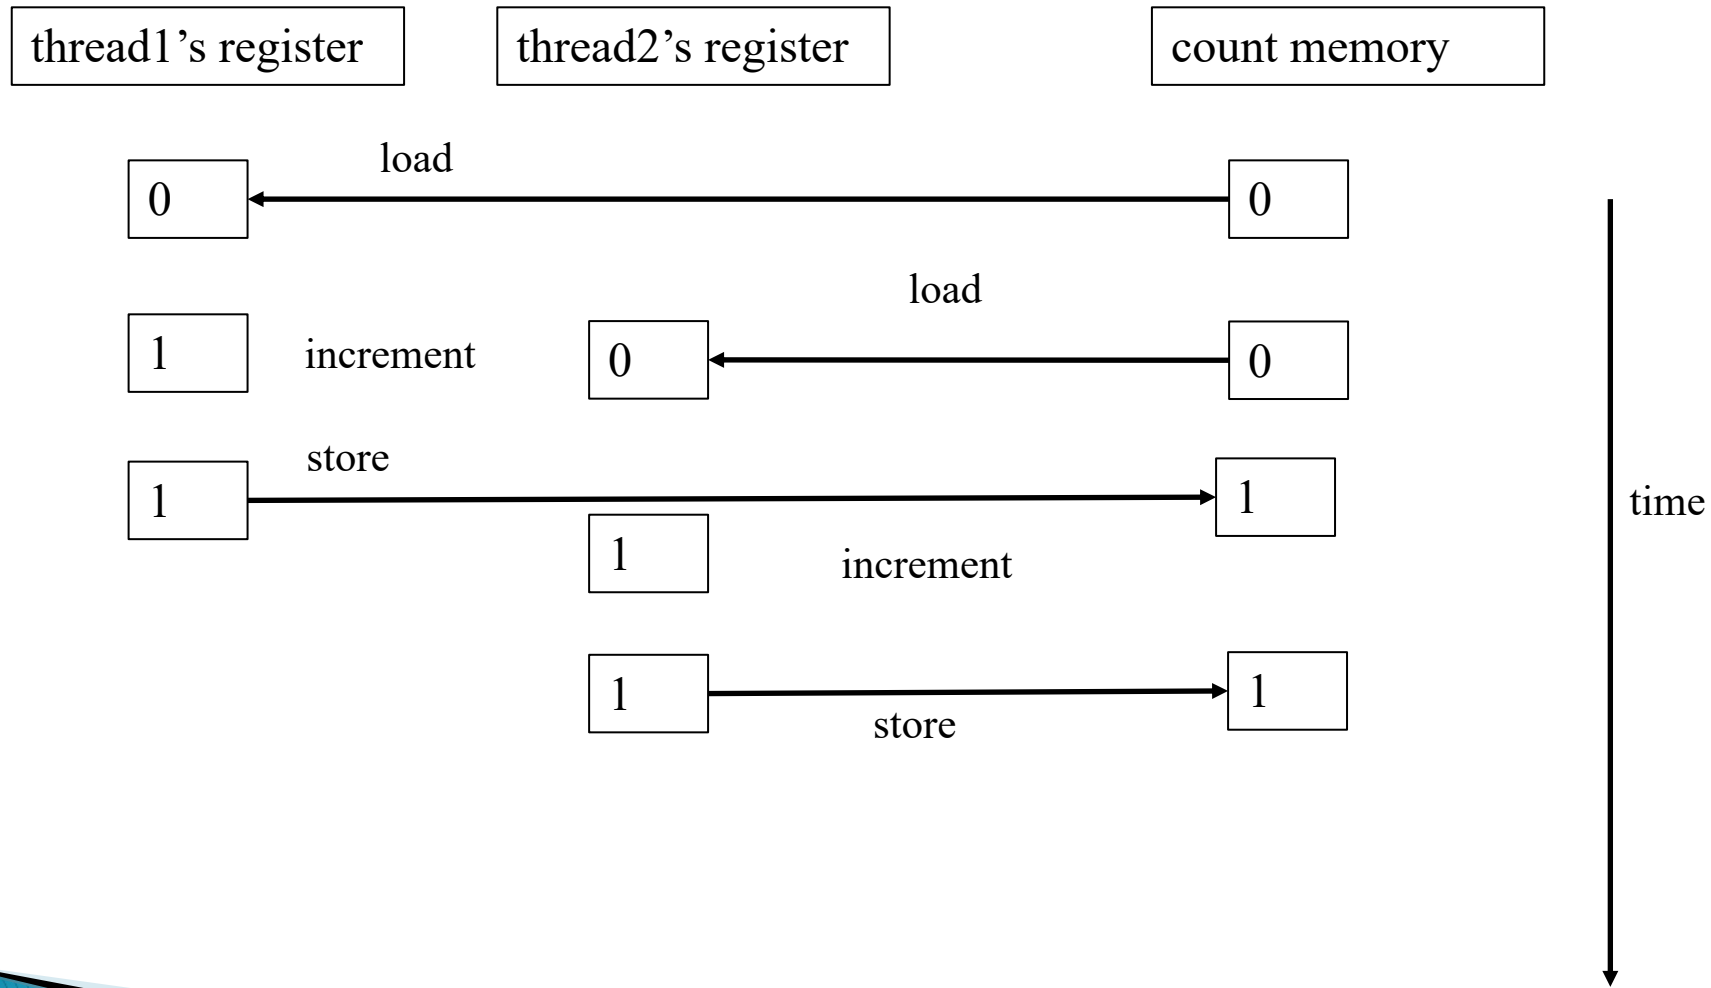
"Normal" operation



Example, cont.

- Suppose the first thread executes `LOAD count;` then `count` is 0 in main memory *and* in the register
 - Now the second thread runs; it executes `LOAD count;` while the first thread is executing `increment`
 - The first thread writes 1 back to memory
 - Then the second thread writes 1 back to memory
- 

Not so normal operation



Race Condition

- So we did `count++` twice to 0 and got 1, not 2.
 - Basically because both threads started with 0: they weren't serialized, and the second thread used "stale" data
- For larger data items, it only gets worse
- If the explanation doesn't make sense to you, no problem, just remember this:

unprotected shared data can be corrupted by concurrent threads



Example

- UnsafeSimpleData contains one int
- That data is not thread-safe
 - We'll see how to make it thread-safe later

```
public class UnsafeSimpleData { // Wrapper for count
    private int count = 0;
    public void increment() {count++;}
    public int getCount() {return count;}
}
```

Example, cont.

- The `UnsafeCounter` class copies a reference to a `SimpleData` object containing the data
- If there's only one of these, there's no race condition

```
public class UnsafeCounter implements Runnable{
    private UnsafeSimpleData sd;
    public UnsafeCounter(UnsafeSimpleData sd) { this.sd = sd; }

    @Override
    public void run() {
        for (int i=0; i<1000; i++) { sd.increment(); }
        System.out.println("count = " + sd.getCount());
    }
}
```

Example, cont.

- main passes the **same** UnsafeSimpleData reference to two copies of UnsafeCounter – i.e., two threads
- Now we have a race condition

```
UnsafeSimpleData sd = new UnsafeSimpleData();  
Thread counter1 = new Thread( new UnsafeCounter(sd) );  
Thread counter2 = new Thread( new UnsafeCounter(sd) );  
  
counter1.start();    counter2.start();  
counter1.join();    counter2.join();  
System.out.println( sd.getCount() );
```



both UnsafeCounter objects use sd
- i.e., shared data


More on the example

- Both copies of `UnsafeCounter` point to the same data on the heap – that's shared data. Pay attention to how that's done
- If they executed concurrently (as a thread's code), then there's a race condition on that data

Race Conditions

- In practice, race conditions are easy to overlook
- They are also very hard to debug – you can test, test, test your app, everything looks good, then you deliver the app to the customer, and the problem shows up
 - So you have to be proactive about finding race conditions

Data Consistency Techniques

- The `count++` example is a very low-level view, but that's where the problem is
 - The solution is at a much higher level: lock the data
1. Make getters/setters **synchronized** (and possibly other methods)
 - Implicitly locks on entry, unlocks on exit
 2. Use thread-safe data structures (that use synchronized methods or locks)
 3. Use explicit locks
- 


Technique #1: *synchronized* Code Blocks

- The **synchronized** keyword can be applied to methods or to blocks of code (curly-bracketed)
- But not to an entire class

```
public synchronized void computeThisThing( ) { ... }
```

or, alternatively:

```
synchronized (object) {  
    ... some code ...  
}
```

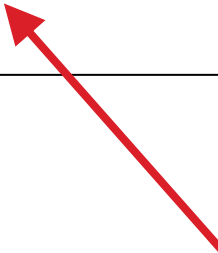


Note: **must** be an object that
is **shared** by the threads

Example: Synchronized

- The SafeSimpleData class protects its data by adding **synchronized** to its methods (not the data)

```
public class SafeSimpleData { // Wrapper for count
    private int count = 0;
    public synchronized void increment() {count++;}
    public synchronized int getCount() {return count;}
}
```



Methods that work on the data are synchronized
EXCEPT for constructors

Example, cont.

- Set up the thread class the same way: it uses the data class

```
public class SafeCounter implements Runnable{
    private SafeSimpleData sd;
    public SafeCounter(SafeSimpleData sd) { this.sd = sd; }

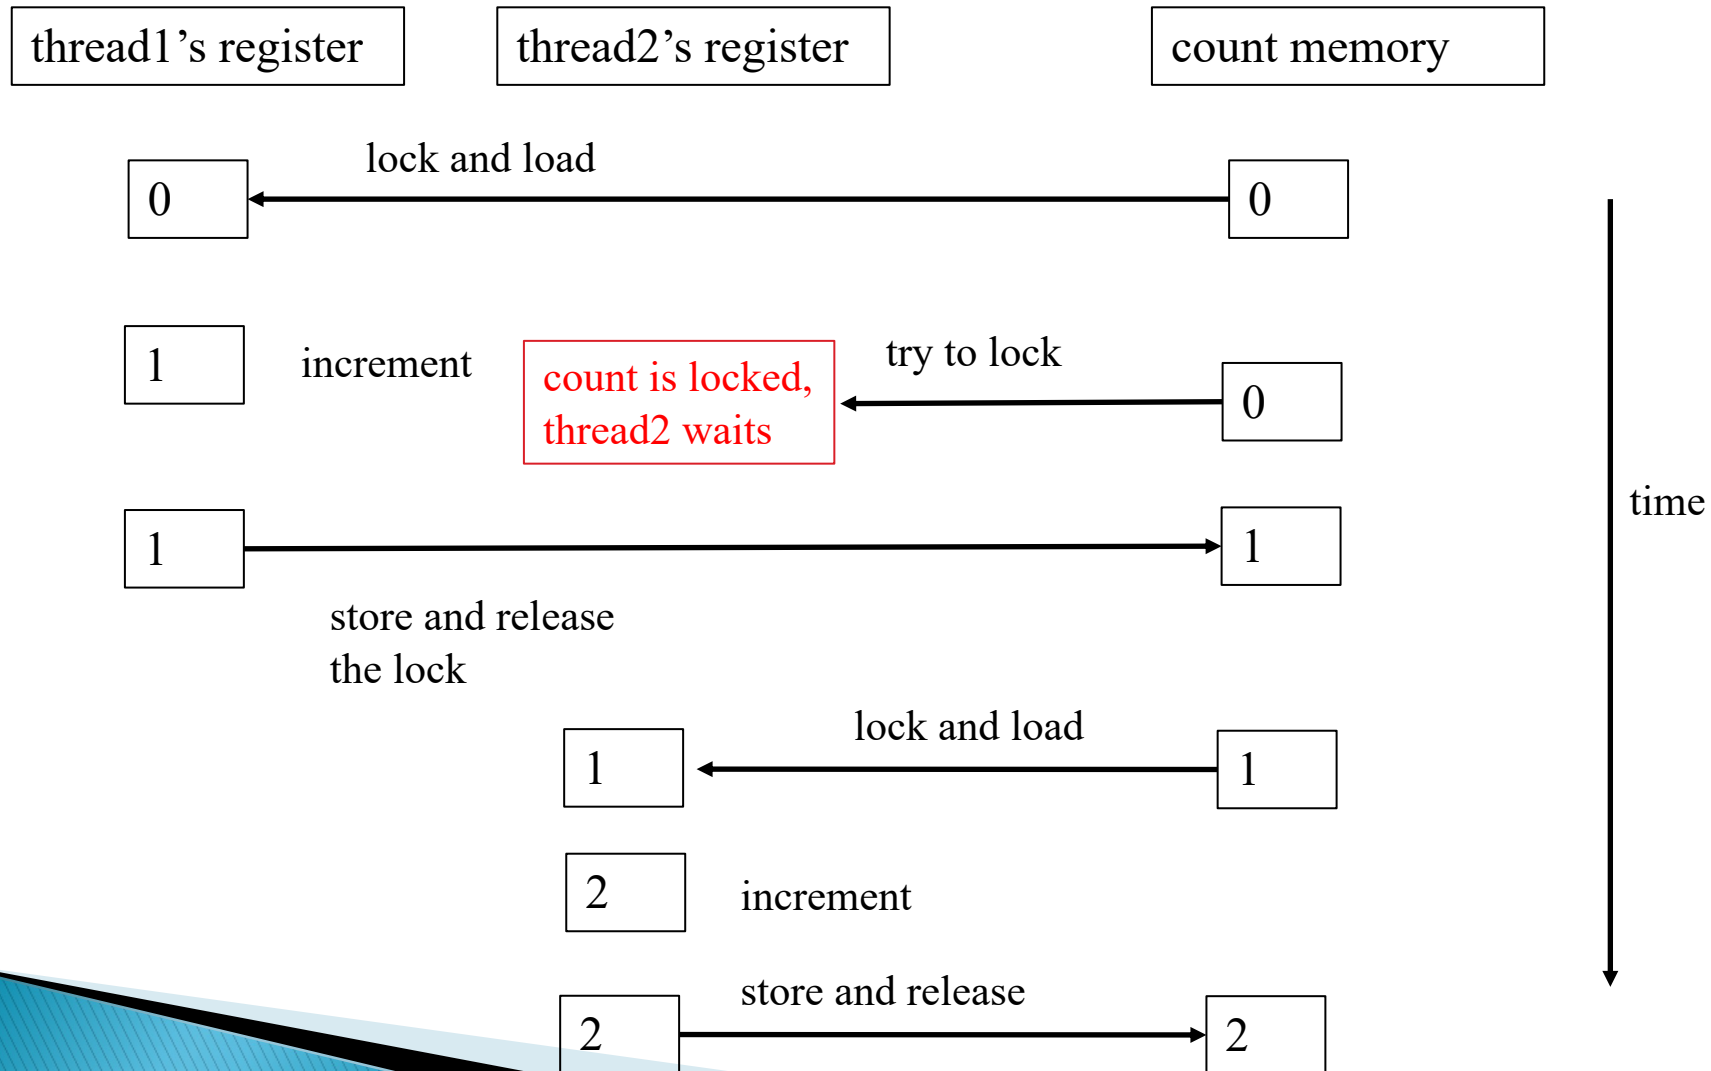
    @Override
    public void run() {
        for (int i=0; i<1000; i++) { sd.increment(); }
        System.out.println("count = " + sd.getCount());
    }
}
```

Example, cont.

- Now the threads take turns
 - each `SafeSimpleData` object has a *lock*, initially open
 - the first thread to access one of the methods *acquires* the lock – i.e. locks the code
 - if a second thread tries to access the object, it waits until first thread *releases* the lock

```
SafeSimpleData sd = new SafeSimpleData();  
Thread counter1 = new Thread( new SafeCounter(sd) );  
Thread counter2 = new Thread( new SafeCounter(sd) );  
counter1.start();    counter2.start();  
counter1.join(); counter2.join();  
System.out.println( sd.getCount() );
```

Synchronized Data



More on the example

- Both copies of `SafeCounter` point to the same data on the heap – that's shared data, just like before
- If they executed concurrently (as a thread's code), there's no race condition, because we synchronized access to the data
 - this *serializes* access to the data

More, cont.

- Note that this does ***not*** imply taking turns - the first thread might run a loop many times before the second one gets a chance, or vice versa – we have no control over the scheduler, we only have control over the lock on the data
- Alternate version: synchronize on something else

```
public void increment() {  
    synchronized(this) {count++;}  
}
```

Technique #2: Thread Safe Data Structures

- Some Java data structures have built-in synchronization
- These are called *thread safe*, because usage by multiple threads will not corrupt their data
- For example, `ArrayList<T>` is not thread safe
 - But if your program doesn't use threads, or the `List` is not shared, use `ArrayList` to avoid the overhead of locking – synchronization is not free
 - Check the docs to be sure

This is not thread-safe

- Example: Don't do this, it's not safe

```
public static void main( ) {  
    ArrayList<String> mylist = new ArrayList<>();  
    ... fill mylist with data ...  
    Thread t1 = new Thread( new MyTask(mylist) );  
    Thread t2 = new Thread( new MyTask(mylist) );  
    t1.start( );  
    t2.start( );  
}
```



both get the same unsafe object

Thread Safe Data Structures

- `java.util.concurrent` contains thread safe data structures like `ArrayBlockingQueue` and `ConcurrentHashMap`
- Basic Collection classes are ***not*** thread safe: `ArrayList`, `HashMap`, ...
- But they can be made thread safe by ***wrapping*** them:

```
List<E> safeArrayList =  
Collections.synchronizedList(new ArrayList<E>());
```

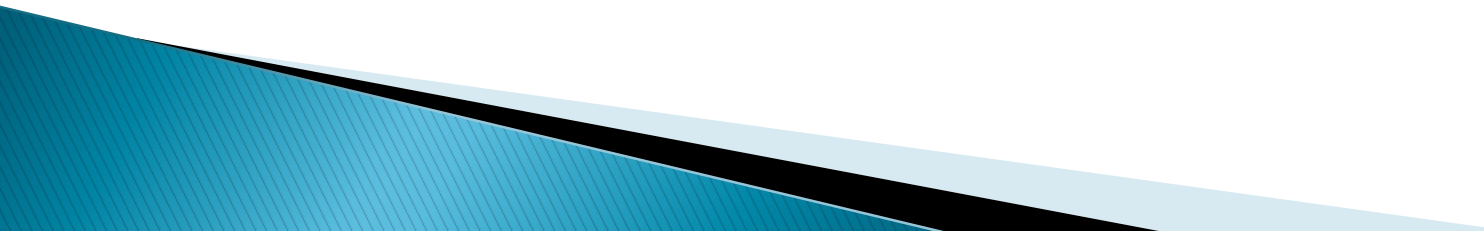


wrapping it to make it thread-safe

Thread Safe Data Structures, cont.

- `Vector` is a thread safe data structure similar to `ArrayList`
- So it can be used when shared between threads

"As of the Java 2 platform v1.2, this class was retrofitted to implement the `List` interface, making it a member of the Java Collections Framework. Unlike the new collection implementations, `Vector` is synchronized. If a thread-safe implementation is not needed, it is recommended to use `ArrayList` in place of `Vector`." – <https://docs.oracle.com/javase/8/docs/api/java/util/Vector.html>



Example

- `ArrayBlockingQueue<>` is thread-safe and ***bounded*** – here, a maximum of 5 elements can be queue.
- `LinkedBlockingQueue<>` is thread-safe and ***unbounded***
- Example uses `BlockingQueue` as shared data

```
private BlockingQueue<Integer> blockingQueue = new
    ArrayBlockingQueue<Integer>(5); ← set to 5 items

Thread producer = new Thread( new NumberMaker( blockingQueue ) );
Thread consumer = new Thread( new QueueProcessor( blockingQueue ) );

producer.start(); consumer.start();
producer.join(); consumer.join();
```

Example, cont.

```
public class NumberMaker implements Runnable {
    private BlockingQueue blockingQueue;

    public NumberMaker(BlockingQueue blockingQueue) {
        this.blockingQueue = blockingQueue;
    }

    @Override
    public void run() {
        for (int i = 0; i < 20; i++ ) {                // put 20 Integer's
            try {
                blockingQueue.put(new Integer(i));
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}
```


Example, cont.

```
public class QueueProcessor implements Runnable {
    private BlockingQueue blockingQueue;

    public QueueProcessor(BlockingQueue blockingQueue) {
        this.blockingQueue = blockingQueue;
    }

    @Override
    public void run() {
        Integer i = null;
        try {
            do {
                i = (Integer)blockingQueue.take(); // take 20 Integer's
                System.out.println("Got " + i);
            } while (i != 20);
        } catch (InterruptedException e) {
            e.printStackTrace();
        }
    }
}
```

More on the example

- The `Runnable` classes use the shared `blockingQueue` to safely move data
- Calls to `put()` and `take()` are synchronized, as are other methods
- But watch out: if you make consecutive method calls, each **individual** call is synchronized, but the sequence of calls is not – i.e., a race condition
 - If you need to do that, synchronize it some other way

```
// This is bad:
```

```
if ( !blockingQueue.contains(thing) ) {  
    blockingQueue.add(thing);    // Race condition  
}
```

Example

- The **AtomicInteger** class is also thread safe
- It's already a class, so it doesn't need any more packaging
- Its methods are not straightforward ...
 - for example, `incrementAndGet()` adds one and returns the new value, instead of just incrementing

```
AtomicInteger sd = new AtomicInteger(); // Don't need an enclosing class
```

```
Thread counter1 = new Thread( new SafeCounter(sd)); // Need to change SafeCounter  
Thread counter2 = new Thread( new SafeCounter(sd)); // to handle AtomicInteger  
counter1.start(); counter2.start();  
counter1.join(); counter2.join();  
System.out.println(sd.get() );
```

volatile

- The **volatile** keyword is related to thread safety, but ...
- ... it does ***NOT*** make a data item thread-safe
- It tells the JVM to automatically write the variable back to memory when it changes value – i.e., it's about visibility, not races – so there's a race condition here
- Only use it for "small" data – like a boolean flag

```
static volatile boolean done = false;
```

```
... then in some thread code:
```

```
if (done) { ... finish up ...}
```

Technique #3: Explicit Locks

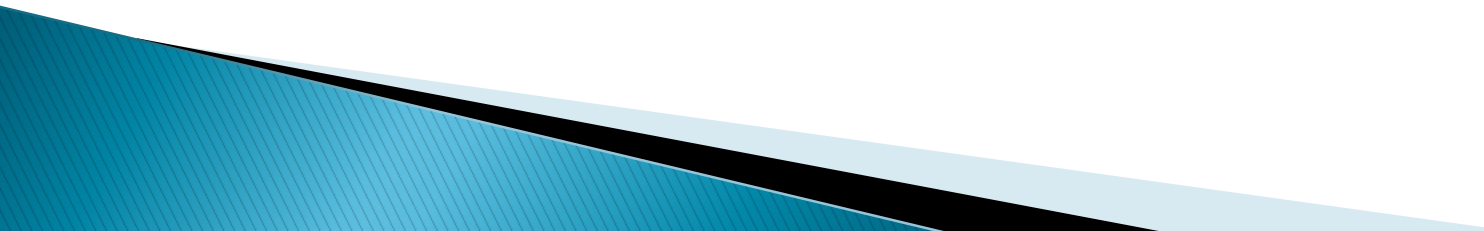
- Java has several kinds of locks; the basic one is **ReentrantLock**.
 - Among its methods: `lock()` and `unlock()`
- So instead of declaring `SafeSimpleData.increment()` as `synchronized`, add a `ReentrantLock` as member data and lock it yourself:

```
private ReentrantLock mylock = new ReentrantLock();  
mylock.lock();  
    // code here  
mylock.unlock();
```



Note: lock must be shared data, too

```
import java.util.concurrent.locks.*;
public class SafeSimpleData { // Alternate implementation
    private int count = 0;
    private ReentrantLock mylock = new ReentrantLock();
    public void increment() {
        mylock.lock();
        try { count++; } // try block is required
        finally {
            mylock.unlock(); // Always unlock in finally
        }
    }
    // similarly for getCount() ...
}
```



More on the example

- This version is more awkward – lower level programming – than using `synchronized`
- It is used when you only need to lock part of your code, not an entire method
 - Lock the part that uses shared data – called the *critical region* – and leave any safe parts unlocked. This is a (small) optimization: the less stuff you lock, the more chance for parallelism, because you're not blocking threads from running

GUI Threading

- A JavaFX application runs on a thread started by `launch()`
- If you have any tasks that take a long time, the GUI will be unresponsive until that task completes
 - Makes users unhappy
- Instead, you can create another thread for the task to run in
- This keeps the GUI responsive – the main UI thread will respond to button clicks, etc.
- Writing your own task class is possible, but there's a built-in threading mechanism called ***Task<T>***
 - It implements `Runnable`

GUI Threading, cont.

- `Task<>` still needs to be put in a `Thread`, then call `start()`, which invokes `call()`
- When the `Task` is finished, it sends a `WorkerStateEvent` that can be handled by `Task.setOnSucceeded()` by overriding its `handle()` method
 - The pattern here is common, but the details are JavaFX-specific: other GUIs will have other ways to do all this
- This part is done on the UI thread after the `Task` has finished

```
drawHamburg.setOnAction(new EventHandler<ActionEvent>() {
    @Override
    public void handle(ActionEvent event) {
        // Task<> is built in to JavaFX, implements Runnable,
        // requires override of call() instead of run()
        Task<Void> task = new Task<Void>() {
            @Override protected Void call() throws Exception {
                // Load the image file
                File file = new File("hamburg.png");
                ... more code to load the image ...
            }

            task.setOnSucceeded(new EventHandler<WorkerStateEvent>() {
                @Override
                public void handle(WorkerStateEvent event) {
                    System.out.println("in setOnSucceeded()");
                    imageView.setImage(image);
                }
            });
        });
        Thread t = new Thread(task);
        t.start();
    }
});
```

More on the example

- This is a simple example – the file is local and loading it doesn't take long, so a long operation is simulated with `sleep()`
- The other GUI buttons remain responsive during the operation
- The non-threaded version stalls the GUI until it finishes
- Other GUI's have different versions of the same thing - for example, Android has `AsyncTask<>`